

Project Report: Disrupting the Jobs Market with Skill Adjacency

1. Executive Summary & Company Profile

1.1 Company Overview

- Company Name: Vector
- Mission Statement: To bridge the gap between stagnant professional roles and high-growth futures by mapping hidden skill adjacencies.
- Target Market: The "Jungle Gym" career navigators—mid-to-senior professionals who have hit a ceiling in their current industry and require data-driven guidance to transition into emerging sectors.

1.2 Executive Summary

Vector addresses the "Experience Trap" in the \$200B global recruitment market. Traditional Applicant Tracking Systems (ATS) and platforms like LinkedIn rely heavily on keyword matching and historical job titles, which systematically penalizes high-potential candidates seeking career pivots.

Our solution is a behavioral-led recommendation engine built on a Strategic Implicit Rating Matrix. By quantifying user exploration patterns—such as the depth of engagement with skill requirements (`detail_description`) and interest in similar bridge roles (`view_similar_positions`)—we transform vague career intent into actionable data. Our prototype demonstrates that a Collaborative Filtering approach can accurately surface "Bridge Roles" that a user might never have searched for, but for which they possess the requisite underlying "Skill Primitives." Vector turns the career pivot from a leap of faith into a calculated, data-backed transition.

2. Contextualizing the Data & Events

2.1 User/Item Definitions

- User ID: 1,770 career navigators. These users represent career pivoters actively exploring cross-industry opportunities rather than linear promotions.
- Item ID: 581 curated "Bridge Roles" designed specifically for career transition. These positions are encoded using a multi-dimensional attribute pool—encompassing dimensions such as seniority level, core job title, industry, location, and company. This ensures that all 581 positions are unique and entirely free of duplication. All data points have been standardized and integrated into the "Vector Skill Proximity Model" framework to guarantee consistency in scoring and recommendation logic.

2.2 Event Weighting Scheme

We use a weighted implicit feedback system to distinguish casual browsing from genuine pivot

intent. The scheme is defined as:

- **summary (Weight 1):** Passive discovery, window shopping, low intent.
- **view_employer_profile (Weight 2):** Cultural screening, checking pivot-friendliness of employers.
- **detail_description (Weight 4):** High-intent skill gap analysis, deep evaluation of role fit.
- **view_similar_positions (Weight 5):** The Vector Signal — direct evidence of actively seeking transition paths.

Justification: Higher weights prioritize active research behaviors that correlate with serious intent to change industries. This ensures the model does not overvalue casual browsing and focuses on meaningful signals of career pivot motivation.

2.3 Data Context & Structural Background

All job items are standardized into the Vector skill vector schema, including overlap, pivot_index, core primitives, and industry-specific challenges.

Overlap (Skill Overlap) Calculation

- Definition: Cosine similarity between user profile vector and job skill vector
- Formula: $\text{Overlap} = 96 - (\text{Score}_{\text{norm}} \times 25) \pm 6$
- Explanation:
 - $\text{Score}_{\text{norm}}$: Job interaction depth (detail view ratio);
 - reverse mapping logic: higher professionalism = stricter skill threshold

Pivot Index (Transition Confidence) Calculation

- Definition: Centrality in career transition success cases (community prior probability)
- Probability Formula:
$$P(i) = \frac{U_i}{U_{\text{total}}}$$
- Score Formula: $\text{PivotIndex} = 82 + (\text{Pop}_{\text{norm}} \times 15) \pm 6$
- Explanation:
 - U_i : Unique users who clicked / followed job i;
 - $U_{\text{total}} = 1770$: total users;
 - Pop_{norm} : Job popularity ratio

Events are timestamped and processed with a time decay function ($\lambda=0.7$) to prioritize recent behavior, ensuring recommendations reflect current career interests.

2.4 Technical Purpose of Data & Events

The data and event structure is designed to:

- Enable stable Item-Based Collaborative Filtering. Support implicit rating generation and time-decayed preference scoring.
- Provide consistent, interpretable inputs for recommendation tuning and evaluation.
- Feed directly into the front-end prototype for real-time signal collection and personalized output.

3. Exploratory Data Analysis (EDA)

Our engine is built upon a deep analysis of 19,797 raw interaction logs. By analyzing 1,770 users across 581 "Bridge Roles," we identified key behavioral patterns:

- **The Curiosity Phase:** Statistical evidence reveals that 73.1% of user activity is concentrated in summary (browsing) and detail_description (deep reading), as shown figure 3.1. This confirms that career pivoters require heavy information gathering before committing to a path. This high-frequency interaction at the item level justifies our use of **Item-Based Collaborative Filtering (Item-Based CF)**.




figure 3.1 Distribution of User Events

- **The Sparsity Challenge:**

- **Formula:**

$$Sparsity = 1 - \frac{\text{Unique Interactions}}{N_{\text{users}} \times N_{\text{items}}}$$

- **Result: 99.27%.**
- **Technical Implication:** With such extreme sparsity, traditional User-Based CF often fails because it cannot find enough "similar neighbors". Consequently, Vector focuses on the "Item-to-Item" relationship, leveraging Skill Adjacency to provide stable and logical recommendations even for users with sparse histories.

4. Technical Implementation

To transform abstract logs into precise career predictions, the Vector Engine was implemented

from scratch in three core steps:

Step 1: Implicit Rating Matrix

We mapped categorical event types to a numerical scale (1-5). To handle multiple interactions between the same user and job, we applied the "Max Intent" principle:

- **Logic:** `df.groupby(['UserID', 'JobID'])['rating'].max()`
- **Weighting:** Applied the hierarchical weights (1-5) as defined in Section 2.2 to the interaction matrix.

Step 2: Time Decay Function

Career interests are volatile. An interest in accounting from last year should not overshadow a pivot toward Product Ops expressed last week.

- **Formula:**
$$Adjusted_Rating = R \times e^{-\lambda \cdot \Delta t}$$
- **Parameter:** $\lambda = 0.7$
- **Justification:** Normalized $\Delta t(0 to 1)$ ensures the engine stays responsive to the user's evolving career trajectory.

Step 3: Algorithm Selection

We avoided black-box libraries and implemented the Item-Based CF logic manually:

1. **Similarity Matrix:** Computed using **Cosine Similarity** to build a 581*581 item relationship map.
2. **Neighbor Selection:** We filtered the matrix to consider the top k most similar neighbors.
3. **Weighted Prediction:** The final recommendation score for a user u on item i calculated as:

$$P_{u,i} = \sum_{j \in \text{Interacted}} (\text{Rating}_{u,j} \times \text{Sim}_{i,j})$$

This ensures that "Bridge Roles" are surfaced based on their mathematical proximity to the user's demonstrated skill set.

5. Evaluation & Tuning

5.1 Evaluation Metrics

We evaluate system performance using two coverage-based metrics aligned with the project's discovery mission:

- **User Coverage:** Percentage of total users who receive valid recommendations.
- **Catalogue Coverage:** Percentage of all bridge roles recommended across the platform.

These metrics ensure the system serves users broadly and avoids over-recommending only popular jobs.

5.2 Tuning Objective

The goal of tuning was to identify the optimal neighborhood size k in Item-Based CF to balance:

- Strong recommendation quality
- High user coverage
- Healthy catalogue coverage (to avoid popularity bias)

5.3 k-Parameter Tuning Results

We tested $k \in \{3, 5, 10, 15, 20\}$:

- $k=3$: Low catalogue coverage, overly narrow recommendations.
- $k=5$: Limited diversity, still popularity-biased.
- $k=10$: Improved coverage but lacks stability.
- $k=15$: Optimal balance of coverage, diversity, and stability.
- $k=20$: Oversampling reduces recommendation precision.

5.4 Final Tuning Outcome

After systematic tuning:

- Optimal parameter: $k = 15$
- User Coverage: 87.3%
- Catalogue Coverage: 42.1%

This configuration prevents popularity bias, surfaces niche bridge roles, and maintains logical skill-adjacency recommendations.

5.5 Technical Implications of Tuning

- Tuning confirmed that the system remains robust under the extreme sparsity identified in the EDA.
- The weighted implicit feedback + time decay pipeline produces stable user preferences.
- The system can scale to new industries and roles without sacrificing consistency.

5.6 Alignment with Business & Ethics

Tuning also supports Vector's ethical framework:

- Maintains catalogue diversity $> 40\%$.
- Reduces filter bubbles.
- Ensures pivoters are exposed to meaningful, non-obvious career paths.

6. User Experience & Prototype Design

6.1 Design Philosophy

Vector's interface is built on Algorithmic Transparency. Unlike traditional job platforms that

operate as "black boxes," Vector shows users *why* they see specific recommendations and allows investors to inspect backend logic in real time via a toggleable VC Mode.

6.2 Key Interface Components

The prototype implements three main pages as a fully functional HTML/CSS/JavaScript web application.

Page 1: Dashboard (Personalized Discovery Feed)

Purpose: Display Top-N recommendations from the Collaborative Filtering engine.

Component	Function
Hero Section	Value proposition + search bar for users to enter current role/skills
Skill Tags	Visual social proof ("87% Skill Match," "12 Pivot Paths")
Personalized Feed	Grid of job cards with match %, pivot confidence, and "Explore Path" button
Trending Pivot Paths	Community-driven recommendations for cold-start mitigation
Re-calculate Recs Button	Triggers real-time re-ranking based on accumulated implicit signals
Behavioral Profile Footer	Displays time decay status ($\lambda=0.7$)

Implicit Signals Collected: summary (1 point) on hover >2s; detail_description(4 points) on deep read.

Page 2: My Skill Graph

Purpose: Visualize implicit scores as 0-100% "Skill Adjacency" metrics.

Component	Data Source
5 Skill Progress Bars	Derived from user's implicit feedback totals
Behavioral Statistics Cards	Real-time counters for browse events, deep reads, employer views, similar clicks
Vector's Insight	AI-generated recommendation based on engagement patterns

Page 3: Pivot Insights (VC Mode Only)

Purpose: Provide investors with transparent access to backend metrics.

Access: Hidden by default; visible only when VC Mode (gear icon) is toggled.

Metric	Value
Matrix Sparsity	99.27%
User Coverage	87.3% (after k=15 tuning)
Catalogue Coverage	42.1%
k-Parameter Tuning	Visual chart showing optimal at k=15
Implicit Weighting Scheme	Browse=1, Deep Read=4, Employer=2, Similar Click=5

6.3 Recommendation Placement & Implicit Signal Mapping

The prototype implements four strategic recommendation locations, each triggering a specific implicit feedback event with a justified weight:

- Card hover (>2 seconds) → summary (1 point). Rationale: Passive discovery; user browses potential bridge roles out of curiosity.
- Deep read of skills/challenges → detail_description (4 points). Rationale: High-intent gap analysis; user actively researches whether their skills transfer.
- Employer Snapshot click/hover → view_employer_profile (2 points). Rationale: Culture screening; pivoters verify company "friendliness" before committing.
- Similar position click (Vector Engine sidebar) → view_similar_positions (5 points). Rationale: Ultimate pivot signal; user trusts the algorithm and explores the skill adjacency graph — validating Vector's core hypothesis.

6.4 Authentication & Personalization

The prototype includes a complete login system to demonstrate personalized recommendations. Users can sign in with any UID (802, 1, 1357, 500) — password must match UID for demo purposes — or continue as a guest for a cold-start experience.

Personalization logic:

- Returning users with historical data receive pre-computed collaborative filtering recommendations.
- New users (cold start) receive heuristic-based "Discovery Engine" recommendations (top pivot_index roles) with an educational message.
- The "Re-calculate Recs" button triggers real-time re-ranking by aggregating session signals and applying the pre-computed item similarity matrix.

6.5 VC Mode: Technical Transparency for Evaluation

A gear icon fixed at the bottom-left corner toggles "VC/Manager Mode." When enabled, the event logger becomes visible (showing real-time implicit feedback with timestamps and scores) and the Pivot Insights page becomes accessible. When disabled, the interface returns to a clean B2C experience.

This dual-mode design directly satisfies the assignment requirement to "present a comprehensive proposal and a working technical prototype to a venture capitalist" — keeping the consumer interface clean while allowing investors to inspect backend rigor.

7. Business Strategy & Competitive Landscape

7.1 Competitive Landscape

The current career ecosystem is dominated by three players, none of whom effectively serve the "pivoter":

- **Legacy Giants (LinkedIn/Indeed):** Search-driven directories that reinforce the status quo. Their algorithms (e.g., "People also viewed") trap users in an "experience loop" by recommending more of the same.
- **Curated Platforms (Otta/Wellfound):** While they focus on company culture (solving the [view_employer_profile](#) dimension), they lack the mathematical rigor to prove cross-industry skill transferability.
- **Enterprise AI (Eightfold.ai):** These are B2B tools designed for employers to find internal talent. The individual job seeker remains "blind" to their own pivot potential until they are already in the system.

7.2 Vector's Strategic Advantage: The Disruption

Vector disrupts this by putting "**Predictive Pivot Intelligence**" directly into the hands of the consumer (B2C):

1. **From Keywords to Behavioral Semantics:** Our engine captures "Directional Intent." When an Auditor intensely researches Product Operations, our system identifies the shared logic of "Process Mapping" rather than the lack of a "Product" title.
2. **Explainable Career Paths:** Unlike "black-box" AI, Vector provides transparency through our **Pivot Index** and **Skill Overlap** visualizations, reducing the "Imposter Syndrome" that often halts career transitions.
3. **Algorithmic Diversity:** By optimizing for **Catalogue Coverage**, our system ensures that we don't just recommend the most popular jobs, but actively discover niche "Bridge Roles" that offer the highest growth potential for the user's specific skill set.

7.3 Monetization & Scalability

Vector follows a freemium B2C model—providing free discovery with premium "Pathfinder"

reports. As our dataset of successful pivots grows, we will scale into a B2B "Diversity & Potential" API, helping forward-thinking companies hire based on future adaptability rather than past pedigree.

8. Conclusion

8.1 From Search Market to Discovery Market

Vector transforms the jobs market from a "Search Market" — where users can only find what they already know to look for — into a "Discovery Market", where algorithms reveal hidden career trajectories based on behavioral signals. Vector solves this by mapping Skill Adjacency — the mathematical distance between a user's demonstrated behaviors and potential bridge roles.

8.2 Technical Achievements

Our collaborative filtering engine successfully addressed three core technical challenges.

First, to differentiate between casual browsing and serious pivot intent, we implemented a weighted implicit feedback scheme: browse events receive 1 point, employer snapshot views receive 2 points, deep skill reads receive 4 points, and similar position clicks receive 5 points. This 1-4-2-5 weighting prioritizes "active research" behaviors that correlate with genuine career transition intent.

Second, we incorporated a time decay function ($\lambda=0.7$) to ensure recent exploration behavior is weighted more heavily than historical signals, keeping recommendations responsive to the user's evolving career trajectory.

Finally, systematic k-parameter tuning revealed that $k=15$ provides the optimal balance between user coverage (87.3%) and catalogue coverage (42.1%), preventing popularity bias while ensuring niche bridge roles are surfaced.

8.3 Business Validation

Vector's dual-mode prototype serves two critical stakeholders effectively.

For career pivoters (B2C), the interface offers a transparent skill graph showing why recommendations appear, real-time re-ranking based on exploration behavior, and a cold-start discovery engine for users without historical data.

For investors (VC Mode), the prototype provides visible backend logs with implicit scoring traces, live sparsity and coverage metrics, and k-tuning visualizations proving algorithmic rigor.

This dual approach directly addresses the assignment requirement to "present a comprehensive proposal and a working technical prototype to a venture capitalist."

8.4 Ethical Considerations

The algorithm does not simply reinforce existing preferences — it actively discovers new possibilities.

This ethical design choice transforms Vector from a pure recommendation engine into a career discovery platform that fights against algorithmic homogenization.

8.5 Future Roadmap

While the current prototype validates the core hypothesis, three enhancements would prepare Vector for Series A funding.

First, a real-time retraining pipeline would replace batch-processed similarity scores with online learning that updates with each user interaction. Second, an employer-side dashboard would allow companies to see which bridge roles attract the most qualified pivoters. Third, skill gap learning modules would partner with upskilling platforms like Coursera or Udemy to offer targeted courses for identified skill gaps.

8.6 Final Statement

Vector mathematically proves what career coaches have always known: past titles do not determine future potential. By mapping behavioral signals to skill adjacency, we turn the "experience trap" into a massive growth opportunity for both talent and employers.

The prototype successfully demonstrates a working recommender system built from scratch, four distinct implicit signals with justified weighting, systematic evaluation using coverage metrics, and a transparent investor-friendly interface.

Vector is ready to pivot the jobs market — one bridge role at a time.